

## 5A Optimise Performance, Security and Privacy

### 5A.1 Optimise Performance

The UI/mobile app is the entry point for users to interact with the system, while the backend server is responsible for handling business logic and processing data from both the UI and sensors. The system includes interfaces with third-party services for social media platforms and payment gateways, and the database for storing user data. Deployment is on Amazon Web Services (AWS) which provides tools to manage system load and ensure microservice redundancy. In order to make the best use of the AWS tools, the potential Single Points of Failure (SPoFs) or areas of intensive data processing need to be identified.

#### Single Points of Failure (SPoFs)

The **backend server** hosts the FitRideX Connect system. Although the application is structured as a microservices system which offers more fault tolerance, careful management of the services is essential. Deploying the application on AWS allows use of service management tools but potential SPoFs may still occur. Any issues with the backend server could lead to a complete system outage. For this reason, it is important to configure and manage access to the server correctly.

Service discovery is a mechanism used in microservices architectures to automatically detect and locate services within a network. This is crucial because microservices often run in dynamic environments where service instances can frequently change due to scaling, updates, or failures. AWS defines Availability Zones (AZ), A Service Registry data keeps track of all available service instances and their locations (IP addresses and ports). When a microservice instance starts, it registers with the Service Registry indicating its availability. When a service consumer needs to communicate with another service, it queries the Service Registry to find available instances of the required service and their locations.

Load Balancers are critical components needed to ensure access to microservices hosted on backend servers is maintained. They distribute incoming traffic to multiple servers, ensuring no single server becomes overwhelmed which helps to ensure that the FitRideX application will remain available even if one or more servers fail. By continuously monitoring the health of servers they ensure traffic is only sent to servers that are up and running, and ensure system performance by reducing latency and improved response times by balancing load across multiple servers.

In the AWS framework, server-side discovery is used; a client process makes a request to the Load Balancer which queries the Service Registry, and forwards the request to an available service instance. Using a Load Balancer also allows for session persistence,



which ensures that all requests from a user during a session are sent to the same server. This is important for the FitRideX Connect system which relies on continuous updating of the UI during live and virtual rides.

Availability Zones (AZ) are isolated locations within a cloud provider's region, designed to enhance the availability and reliability of applications and services. Each AZ is a separate data centre with independent power, cooling, and networking infrastructure.

Another potential SPOF is if a Load Balancer fails which can disrupt traffic distribution. To avoid this possible system failure, AWS offers Elastic Load Balancing (ELB) with multiple Load Balancer instances across different AZs to ensure redundancy.

If all instances of a microservice are in a single AZ, an AZ failure can take down the service. Deploying instances across multiple AZs and configuring ELB's auto scaling to handle instance failures and traffic spikes will protect the system against backend server outages.

**Message Brokers** are an important aspect of the FitRideX Connect system as they perform the critical role of inter-service communications. A single message broker instance can fail, disrupting communication between services. By using multiple broker instances deployed in different AZs, FitRideX Connect can ensure message durability and availability.

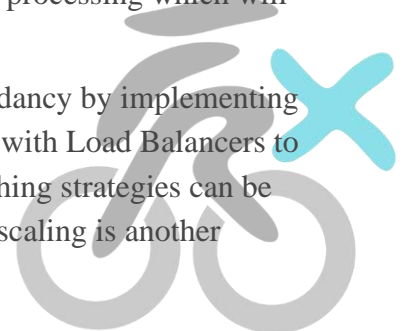
Maintaining **Network Connectivity** is essential for the effective operation of the FitRideX application. Any network issue can result in isolated services. Using AWS Direct Connect and VPN for redundant network paths and ensuring services are deployed in multiple regions can help to mitigate this potential SPoF.

The **Database** is an essential component of the FitRideX Connect system. If there is a single database instance handling all data storage and retrievals, it could be a SPoF. For this reason, it is important to ensure that multiple database instances are deployed. In addition, a database server should be used to balance queries across the multiple database instances and improve performance.

## **Intensive Data Processing**

The FitRideX application processes real-time data from sensors during cycling sessions, which could be resource-intensive, especially during peak usage times such as early mornings and weekends. Additionally, performing complex data analytics to generate training or performance metrics could also be a point of intensive data processing which will also put a higher demand on access to the database.

To address these potential issues, the system will ensure redundancy by implementing multiple server and database instances, and managing system demand with Load Balancers to distribute traffic and processing loads across the multiple servers. Caching strategies can be used to reduce the load on the backend server and database. Dynamic scaling is another



strategy that will be incorporated, which allows for the rapid deployment of additional resources as needed to handle increased data processing demands.

## 5A.2 System Bottlenecks

When designing the FitRideX Connect system, it's crucial to anticipate and address potential performance issues. These issues can arise during high-traffic periods, leading to server overloads and latency. Data spikes, such as those from sudden influxes of user activity or sensor data, can strain databases and processing capabilities. Additionally, resource-intensive operations, like real-time analytics and video streaming, can consume significant computational power and bandwidth, potentially causing bottlenecks and degraded user experiences. Identifying and mitigating these challenges is essential for maintaining a smooth and responsive system.

### High Traffic Periods

High-traffic periods for the FitRideX Connect system are likely to occur during times when users are most active, such as:

- **Early Mornings and Evenings:** Many cyclists prefer to ride before or after work, leading to increased usage during these times.
- **Weekends:** More users may engage in longer rides or group activities on weekends, resulting in higher traffic.
- **Special Events:** Organised cycling events, challenges, or competitions can cause significant spikes in user activity.
- **New Feature Releases:** When new features or updates are rolled out, users may log in simultaneously to explore the new functionalities.

### Data Spikes

Data spikes can be caused by several factors:

- **Simultaneous Data Uploads:** Multiple users uploading ride data at the same time can lead to spikes.
- **Real-Time Analytics:** Resource-intensive operations like real-time performance tracking and analytics can cause temporary data surges.
- **User Growth:** A sudden increase in the number of users, such as after a successful marketing campaign, can lead to higher data volumes.
- **Seasonal Trends:** During peak cycling seasons, such as spring and summer, overall user activity may increase, leading to more data being generated and processed.



## Performance Issues by Component

1. Database
  - Slow query performance due to inadequate indexing or complex joins.
  - Inefficient data retrieval and storage mechanisms.
  - Concurrency issues leading to locking and blocking.
  - High-traffic periods: Increased load can lead to slow response times and timeouts.
  - Data spikes: Sudden influx of data can overwhelm the database's ability to process and store information.
  - Resource-intensive operations: Large-scale data migrations or complex calculations can strain database resources.
2. Server
  - Insufficient CPU or memory to handle concurrent user requests.
  - Inefficient resource allocation, leading to under-utilisation or over-utilisation of server capacity.
  - High-traffic periods: The server may struggle to service requests, leading to increased latency and potential service outages.
  - Data spikes: Sudden increases in data can overwhelm the server's processing capabilities.
  - Resource-intensive operations: Tasks that require significant computational power, such as machine learning or data analytics, can consume server resources.
3. Network
  - Bandwidth constraints causing slow data transfer between components.
  - Network latency due to geographically dispersed components or insufficient infrastructure.
  - High-traffic periods: The network may become congested, leading to delays and packet loss.
  - Data spikes: Sudden increases in data transfer can saturate network links.
  - Resource-intensive operations: Operations that require large amounts of data to be transferred can strain the network.

### 5A.3 System Monitoring Tools

Monitoring and profiling tools are essential for ensuring the optimal performance and reliability of the FitRideX Connect system. These tools can be used to collect, analyse, and visualise system performance metrics and resource utilisation, helping to identify bottlenecks and areas for optimisation. They can be integrated into the FitRideX system allowing the developers and system administrators to proactively address performance issues, enhance user experience, and maintain the overall health of the system. Below is a list of monitoring and profiling tools that can be used to track the performance of the FitRideX Connect system.

1. New Relic: This tool provides real-time insights into application performance, with the ability to monitor and diagnose issues in complex distributed systems.
2. Datadog: Offers a unified platform for infrastructure, application, and log monitoring, allowing for easy tracking of performance metrics and rapid troubleshooting.
3. AppDynamics: Provides application performance monitoring (APM) that helps identify and resolve performance issues in complex applications.
4. Grafana: An open-source analytics and interactive visualisation tool that can be used for monitoring system performance by querying various data sources.
5. Prometheus: An open-source monitoring system and time series database that provides a powerful query language for monitoring and alerting on system performance metrics.
6. JMeter: A Java-based tool designed for load testing and measuring the performance of web applications or FTP and HTTP servers.
7. cAdvisor: A tool designed specifically for container performance monitoring, providing container metrics like CPU, memory, file system, and network usage.
8. ELK Stack (Elasticsearch, Logstash, Kibana): A popular log management and analysis platform that can be used to monitor system performance by analysing log data.
9. Nagios: An open-source computer-software application that monitors systems, networks, and infrastructure. It alerts users when things go wrong and alerts them a second time when the problem has been resolved.
10. Zabbix: An open-source monitoring software tool for diverse IT components, networks, and applications.

## 5A.4 Caching Strategies


For the FitRideX Connect, caching strategies can significantly enhance performance and user experience. Here are some tailored caching strategies:

### Cache-Aside (Lazy Loading)

The system first checks the cache for ride data, user profiles, or training plans. If the data is not found (cache miss), it retrieves the data from the database and then stores it in the cache for future requests. Ideal for frequently accessed data like user profiles and commonly used training plans.

### Write-Through

Data updates, such as new ride statistics or updated training plans, are written to both the cache and the database simultaneously. This ensures that the cache is always up-to-date.



Useful for maintaining consistency between the cache and the database, especially for real-time ride data.

### **Read-Through**

The system interacts only with the cache. If the data is not in the cache, the cache itself retrieves the data from the database and then returns it to the application. Simplifies the application code and ensures that frequently accessed data like leaderboard statistics are quickly available.

### **Write-Behind (Write-Back)**

Data is written to the cache first and then asynchronously written to the database. This can improve write performance but may lead to data loss if the cache fails before the data is written to the database. Suitable for non-critical data where write performance is more important than immediate consistency, such as temporary ride metrics.

### **Distributed Caching**

The cache is distributed across multiple nodes, which helps in scaling the cache horizontally and provides high availability. Essential for handling large volumes of ride data and user interactions, ensuring that the system remains responsive under heavy load.

### **Proactive Caching (Eager Loading)**

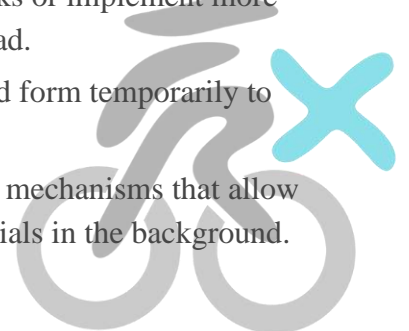
Data is preloaded into the cache based on anticipated requests. This can be done during system start-up or at scheduled intervals. Effective for preloading popular training plans or frequently accessed user data to reduce latency.

## **5A.5 Security and Performance Trade-offs**

Security measures often come with trade-offs that can impact system performance. For instance, encryption and decryption processes can introduce latency, and robust authentication mechanisms can slow down user login times.

By implementing mitigating strategies, security measures can be enhanced while minimising their impact on system performance.

1. **Optimise Encryption:** Use hardware acceleration for encryption tasks or implement more efficient encryption algorithms to reduce the computational overhead.
2. **Caching:** Cache frequently accessed encrypted data in its decrypted form temporarily to reduce the number of encryption and decryption operations.
3. **Asynchronous Authentication:** Using asynchronous authentication mechanisms that allow users to perform other tasks while the system verifies their credentials in the background.



4. **Load Balancing:** Distribute the load across multiple servers to handle the computational demands of security measures more effectively.
5. **Efficient Security Protocols:** Use security protocols that are known for their efficiency, such as TLS 1.3, which offers performance improvements over earlier versions.
6. **Regular Maintenance:** Keep security software and systems up to date with the latest patches and updates to ensure optimal performance.
7. **Profiling and Monitoring:** Continuously monitor and profile security-related processes to identify bottlenecks and optimise them accordingly.

## 5B Security and Privacy Audit

### 5B.1 Personal or Sensitive Data

To perform a comprehensive security and privacy audit for FitRideX Connect, it is crucial to first identify the types of personal or sensitive data that the system handles, including personal identifiable information (PII), health records, financial data, data that could be considered personal or sensitive, and any other information that requires protection under privacy laws. A comprehensive inventory of all data collected, processed, stored, and transmitted by the system is tabled below. The data is also classified based on the required security level, and the risks associated with access breaches to the data are identified.

Description	Security Level	Risks
<b>Personal Identifiable Information (PII)</b>		
User Data: <ul style="list-style-type: none"><li>• Name, age, gender, contact details.</li></ul> Account Information: <ul style="list-style-type: none"><li>• Username, password, profile picture.</li></ul>	High – Requires encryption, access controls, and regular audits.	Unauthorised access to PII can lead to identity theft, phishing attacks, and privacy violations. For FitRideX, this can result in legal penalties, loss of customer trust, and reputational damage.
<b>Health and Training Records</b>		
User Fitness Data: <ul style="list-style-type: none"><li>• Includes health metrics like heart rate, weight, and other fitness-related measurements.</li></ul>	High – Requires encryption, strict access controls, and compliance with health data regulations.	Breaches can expose sensitive health information, leading to privacy violations and potential misuse of health data. FitRideX may face legal consequences and loss of user trust.

